

The background is a deep blue gradient with a starry space texture. On the left side, there are several overlapping circular patterns. Some are solid white lines, while others are dashed. A large circular scale with numerical markings from 140 to 260 in increments of 10 is visible. Small white arrows are scattered throughout the circular patterns, some pointing clockwise and others counter-clockwise. The overall aesthetic is technical and futuristic.

# EDA & QA TORNADOS

TOMAS TORTOSA MOLTO

head()

```
tornado.head()
```

Python

	om	yr	mo	dy	date	time	tz	datetime_utc	st	stf	...	elon	len	wid	ns	sn	f1	f2	f3	f4	fc
0	192	1950	10	1	1950-10-01	21:00:00	America/Chicago	1950-10-02T03:00:00Z	OK	40	...	-102.3	15.8	10	1	1	25	0	0	0	False
1	193	1950	10	9	1950-10-09	02:15:00	America/Chicago	1950-10-09T08:15:00Z	NC	37	...	0.0	2.0	880	1	1	47	0	0	0	False
2	195	1950	11	20	1950-11-20	02:20:00	America/Chicago	1950-11-20T08:20:00Z	KY	21	...	0.0	0.1	10	1	1	177	0	0	0	False
3	196	1950	11	20	1950-11-20	04:00:00	America/Chicago	1950-11-20T10:00:00Z	KY	21	...	0.0	0.1	10	1	1	209	0	0	0	False
4	197	1950	11	20	1950-11-20	07:30:00	America/Chicago	1950-11-20T13:30:00Z	MS	28	...	0.0	2.0	37	1	1	101	0	0	0	False

5 rows x 27 columns



tail()

```
tornado.tail()
```

Python

	om	yr	mo	dy	date	time	tz	datetime_utc	st	stf	...	elon	len	wid	ns	sn	f1	f2	f3	f4	fc
68688	621922	2022	9	28	2022-09-28	13:32:00	America/Chicago	2022-09-28T18:32:00Z	FL	12	...	-80.8841	3.00	100	1	1	9	0	0	0	False
68689	621923	2022	9	30	2022-09-30	10:25:00	America/Chicago	2022-09-30T15:25:00Z	NC	37	...	-78.3011	0.74	20	1	1	19	0	0	0	False
68690	621924	2022	9	30	2022-09-30	13:22:00	America/Chicago	2022-09-30T18:22:00Z	NC	37	...	-76.7147	0.70	12	1	1	13	0	0	0	False
68691	621900	2022	9	4	2022-09-04	15:44:00	America/Chicago	2022-09-04T20:44:00Z	OH	39	...	-80.6555	0.07	15	1	1	99	0	0	0	False
68692	621901	2022	9	9	2022-09-09	23:21:00	America/Chicago	2022-09-10T04:21:00Z	SC	45	...	-79.7537	0.68	125	1	1	19	0	0	0	False

5 rows x 27 columns

# describe()

```
tornado.describe()
```

Python

	om	yr	mo	dy	stf	mag	inj	fat	loss	slat	...	elat
count	68693.000000	68693.000000	68693.000000	68693.000000	68693.000000	67937.000000	68693.000000	68693.000000	4.152300e+04	68693.000000	...	68693.000000
mean	113201.815542	1991.854061	5.968541	15.930881	29.220255	0.778721	1.418689	0.08931	2.020898e+06	37.129386	...	22.960651
std	226621.993899	19.565158	2.444656	8.750070	15.013273	0.895790	18.114752	1.47212	3.039588e+07	5.099005	...	18.528144
min	1.000000	1950.000000	1.000000	1.000000	1.000000	0.000000	0.000000	0.00000	5.000000e+01	17.721200	...	0.000000
25%	285.000000	1976.000000	4.000000	8.000000	18.000000	0.000000	0.000000	0.00000	1.000000e+04	33.180000	...	0.000000
50%	588.000000	1995.000000	6.000000	16.000000	28.000000	1.000000	0.000000	0.00000	5.000000e+04	37.000000	...	32.550000
75%	1118.000000	2008.000000	7.000000	24.000000	42.000000	1.000000	0.000000	0.00000	5.000000e+05	40.920000	...	38.650000
max	622080.000000	2022.000000	12.000000	31.000000	78.000000	5.000000	1740.000000	158.00000	2.800100e+09	61.020000	...	61.020000

8 rows x 21 columns



# info()

```
tornado.info()
```

Python

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 68693 entries, 0 to 68692
Data columns (total 27 columns):
#   Column          Non-Null Count  Dtype
---  -
0   om              68693 non-null  int64
1   yr              68693 non-null  int64
2   mo              68693 non-null  int64
3   dy              68693 non-null  int64
4   date            68693 non-null  object
5   time            68693 non-null  object
6   tz              68693 non-null  object
7   datetime_utc    68693 non-null  object
8   st              68693 non-null  object
9   stf             68693 non-null  int64
10  mag             67937 non-null  float64
11  inj             68693 non-null  int64
12  fat             68693 non-null  int64
13  loss            41523 non-null  float64
14  slat            68693 non-null  float64
15  slon            68693 non-null  float64
16  elat            68693 non-null  float64
17  elon            68693 non-null  float64
18  len             68693 non-null  float64
19  wid             68693 non-null  int64
...
25  f4              68693 non-null  int64
26  fc              68693 non-null  bool
dtypes: bool(1), float64(7), int64(14), object(5)
memory usage: 13.7+ MB
```

# Shape()

```
tornado.shape
```

Python

```
(68693, 27)
```



Eliminamos columnas que no son necesarias (ID del tornado, zona horaria y DateTime)

```
tornado = tornado.drop('om', axis=1)
tornado = tornado.drop('tz', axis=1)
tornado = tornado.drop('datetime_utc', axis=1)
tornado
```

Python

	yr	mo	dy	date	time	st	stf	mag	inj	fat	...	elon	len	wid	ns	sn	f1	f2	f3	f4	fc
0	1950	10	1	1950-10-01	21:00:00	OK	40	1.0	0	0	...	-102.3000	15.80	10	1	1	25	0	0	0	False
1	1950	10	9	1950-10-09	02:15:00	NC	37	3.0	3	0	...	0.0000	2.00	880	1	1	47	0	0	0	False
2	1950	11	20	1950-11-20	02:20:00	KY	21	2.0	0	0	...	0.0000	0.10	10	1	1	177	0	0	0	False
3	1950	11	20	1950-11-20	04:00:00	KY	21	1.0	0	0	...	0.0000	0.10	10	1	1	209	0	0	0	False
4	1950	11	20	1950-11-20	07:30:00	MS	28	1.0	3	0	...	0.0000	2.00	37	1	1	101	0	0	0	False
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
68688	2022	9	28	2022-09-28	13:32:00	FL	12	0.0	0	0	...	-80.8841	3.00	100	1	1	9	0	0	0	False
68689	2022	9	30	2022-09-30	10:25:00	NC	37	0.0	0	0	...	-78.3011	0.74	20	1	1	19	0	0	0	False
68690	2022	9	30	2022-09-30	13:22:00	NC	37	NaN	0	0	...	-76.7147	0.70	12	1	1	13	0	0	0	False
68691	2022	9	4	2022-09-04	15:44:00	OH	39	0.0	0	0	...	-80.6555	0.07	15	1	1	99	0	0	0	False
68692	2022	9	9	2022-09-09	23:21:00	SC	45	0.0	0	0	...	-79.7537	0.68	125	1	1	19	0	0	0	False

68693 rows x 24 columns

# Aplicamos LabelEncoder a las columnas ST y FC

```
from sklearn import preprocessing
le = preprocessing.LabelEncoder()
tornado.st = le.fit_transform(tornado.st)
tornado.fc = le.fit_transform(tornado.fc)
tornado
```

Python

	yr	mo	dy	date		time	st	stf	mag	inj	fat	...	elon	len	wid	ns	sn	f1	f2	f3	f4	fc
0	1950	10	1	1950-10-01	1900-01-01	21:00:00	36	40	1.0	0	0	...	-102.3000	15.80	10	1	1	25	0	0	0	0
1	1950	10	9	1950-10-09	1900-01-01	02:15:00	27	37	3.0	3	0	...	0.0000	2.00	880	1	1	47	0	0	0	0
2	1950	11	20	1950-11-20	1900-01-01	02:20:00	17	21	2.0	0	0	...	0.0000	0.10	10	1	1	177	0	0	0	0
3	1950	11	20	1950-11-20	1900-01-01	04:00:00	17	21	1.0	0	0	...	0.0000	0.10	10	1	1	209	0	0	0	0
4	1950	11	20	1950-11-20	1900-01-01	07:30:00	25	28	1.0	3	0	...	0.0000	2.00	37	1	1	101	0	0	0	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
68688	2022	9	28	2022-09-28	1900-01-01	13:32:00	9	12	0.0	0	0	...	-80.8841	3.00	100	1	1	9	0	0	0	0
68689	2022	9	30	2022-09-30	1900-01-01	10:25:00	27	37	0.0	0	0	...	-78.3011	0.74	20	1	1	19	0	0	0	0
68690	2022	9	30	2022-09-30	1900-01-01	13:22:00	27	37	NaN	0	0	...	-76.7147	0.70	12	1	1	13	0	0	0	0
68691	2022	9	4	2022-09-04	1900-01-01	15:44:00	35	39	0.0	0	0	...	-80.6555	0.07	15	1	1	99	0	0	0	0
68692	2022	9	9	2022-09-09	1900-01-01	23:21:00	41	45	0.0	0	0	...	-79.7537	0.68	125	1	1	19	0	0	0	0


68693 rows x 24 columns



Muestro los tornados de cada año para pensar si hay que hacer binarización

```
plt.hist(tornado.yr, bins=20, color='blue', alpha=0.7)
plt.xlabel('Año')
plt.ylabel('Tornados')
plt.show()
```

Python



# Aplico binarización de 4 bins

```
bins = panda.qcut(tornado.yr, q=4, labels=False)
bins
```

Python

```
0      0
1      0
2      0
3      0
4      0
..
68688   3
68689   3
68690   3
68691   3
68692   3
Name: yr, Length: 68693, dtype: int64
```

```
bins = panda.qcut(tornado.yr, q=4)
bins
```

Python

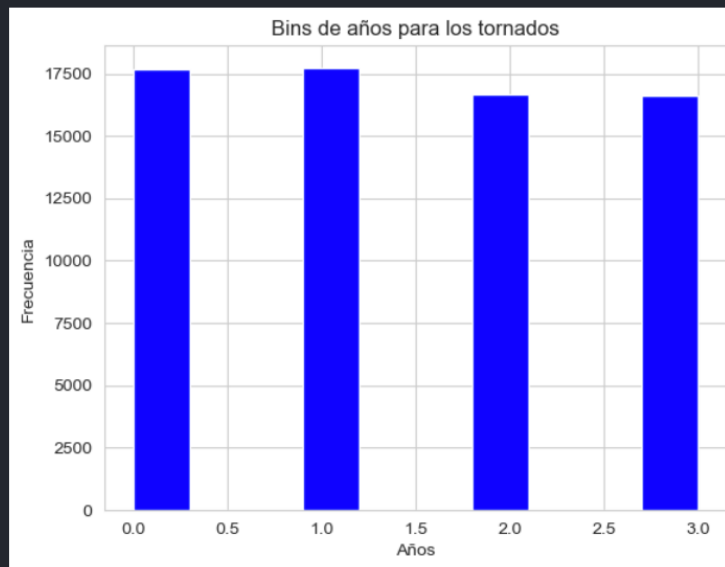
```
0      (1949.999, 1976.0]
1      (1949.999, 1976.0]
2      (1949.999, 1976.0]
3      (1949.999, 1976.0]
4      (1949.999, 1976.0]
...
68688   (2008.0, 2022.0]
68689   (2008.0, 2022.0]
68690   (2008.0, 2022.0]
68691   (2008.0, 2022.0]
68692   (2008.0, 2022.0]
Name: yr, Length: 68693, dtype: category
Categories (4, interval[float64, right]): [(1949.999, 1976.0] < (1976.0, 1995.0] < (1995.0, 2008.0] < (2008.0, 2022.0]]
```



Muestro bins, los años no aparecen bien

```
plt.hist(panda.qcut(tornado.yr, q=4, labels=False), color='blue')  
plt.title('Bins de años para los tornados')  
plt.xlabel('Años')  
plt.ylabel('Frecuencia')  
plt.show()
```

Python



Obtengo el valor de los bins y me quedo con el index

```
years = bins.value_counts().index  
years
```

Python

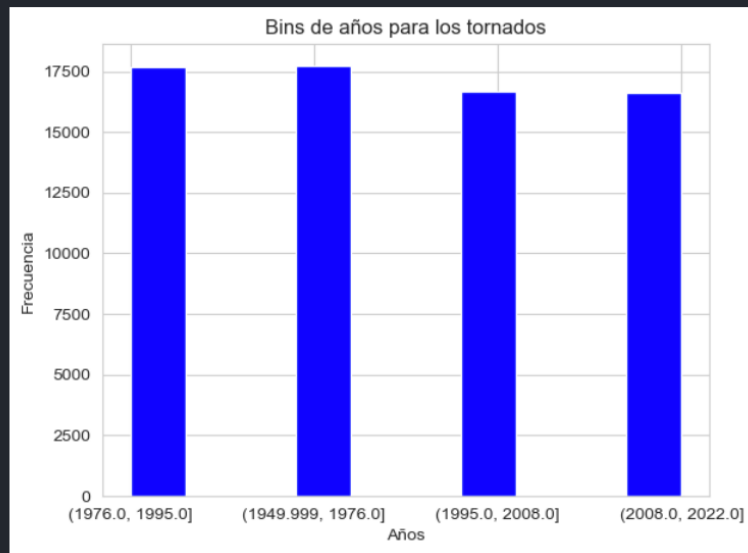
```
CategoricalIndex([(1976.0, 1995.0], (1949.999, 1976.0], (1995.0, 2008.0],  
                 (2008.0, 2022.0]],  
                 categories=[(1949.999, 1976.0], (1976.0, 1995.0], (1995.0, 2008.0], (2008.0, 2022.0]], ordered=True, dtype='category', name='yr')
```



## Muestro gráfico de bins cargando los años anteriores

```
plt.hist(panda.qcut(tornado.yr, q=4, labels=False), color='blue')  
plt.title('Bins de años para los tornados')  
plt.xlabel('Años')  
plt.ylabel('Frecuencia')  
plt.xticks(range(4), years)  
plt.show()
```

Python



# Compruebo columnas nulas

```
tornado.isnull().sum()
```

Python

yr	0
mo	0
dy	0
date	0
time	0
datetime_utc	0
st	0
stf	0
mag	756
inj	0
fat	0
loss	27170
slat	0
slon	0
elat	0
elon	0
len	0
wid	0
ns	0
sn	0
f1	0
f2	0
f3	0
f4	0
fc	0
dtype: int64	



## Calculo media de la columna mag

```
mag = tornado.mag.mean()
print("Esto es media de mag", mag)
magNotNulls = tornado.mag.mean(skipna=True)
print("Esto es media de mag sin nulls", magNotNulls)
```

Python

```
Esto es media de mag 0.778721462531463
Esto es media de mag sin nulls 0.778721462531463
```

Como la media no afecta, elimino filas que no tienen magnitud

```
tornado = tornado.dropna(subset=['mag'])
tornado
```

Python

	yr	mo	dy	date		time	st	stf	mag	inj	fat	...	elon	len	wid	ns	sn	f1	f2	f3	f4	fc
0	1950	10	1	1950-10-01	1900-01-01	21:00:00	36	40	1.0	0	0	...	-102.3000	15.80	10	1	1	25	0	0	0	0
1	1950	10	9	1950-10-09	1900-01-01	02:15:00	27	37	3.0	3	0	...	0.0000	2.00	880	1	1	47	0	0	0	0
2	1950	11	20	1950-11-20	1900-01-01	02:20:00	17	21	2.0	0	0	...	0.0000	0.10	10	1	1	177	0	0	0	0
3	1950	11	20	1950-11-20	1900-01-01	04:00:00	17	21	1.0	0	0	...	0.0000	0.10	10	1	1	209	0	0	0	0
4	1950	11	20	1950-11-20	1900-01-01	07:30:00	25	28	1.0	3	0	...	0.0000	2.00	37	1	1	101	0	0	0	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
68687	2022	9	28	2022-09-28	1900-01-01	03:56:00	9	12	0.0	0	0	...	-80.0680	0.20	50	1	1	99	0	0	0	0
68688	2022	9	28	2022-09-28	1900-01-01	13:32:00	9	12	0.0	0	0	...	-80.8841	3.00	100	1	1	9	0	0	0	0
68689	2022	9	30	2022-09-30	1900-01-01	10:25:00	27	37	0.0	0	0	...	-78.3011	0.74	20	1	1	19	0	0	0	0
68691	2022	9	4	2022-09-04	1900-01-01	15:44:00	35	39	0.0	0	0	...	-80.6555	0.07	15	1	1	99	0	0	0	0
68692	2022	9	9	2022-09-09	1900-01-01	23:21:00	41	45	0.0	0	0	...	-79.7537	0.68	125	1	1	19	0	0	0	0

67937 rows x 24 columns



## Compruebo información para columna loss

```
print("Media de loss:", tornado.loss.median())  
print("Minimo de loss:", tornado.loss.min())  
print("Maximo de loss:", tornado.loss.max())  
print("La desviación típica es:", tornado.loss.std())
```

✓ 0.0s

Python

Media de loss: 50000.0

Minimo de loss: 50.0

Maximo de loss: 2800100000.0

La desviación típica es: 23651659.21147494

Guardo en la columna loss números randoms a partir de la desviación típica

```
import math

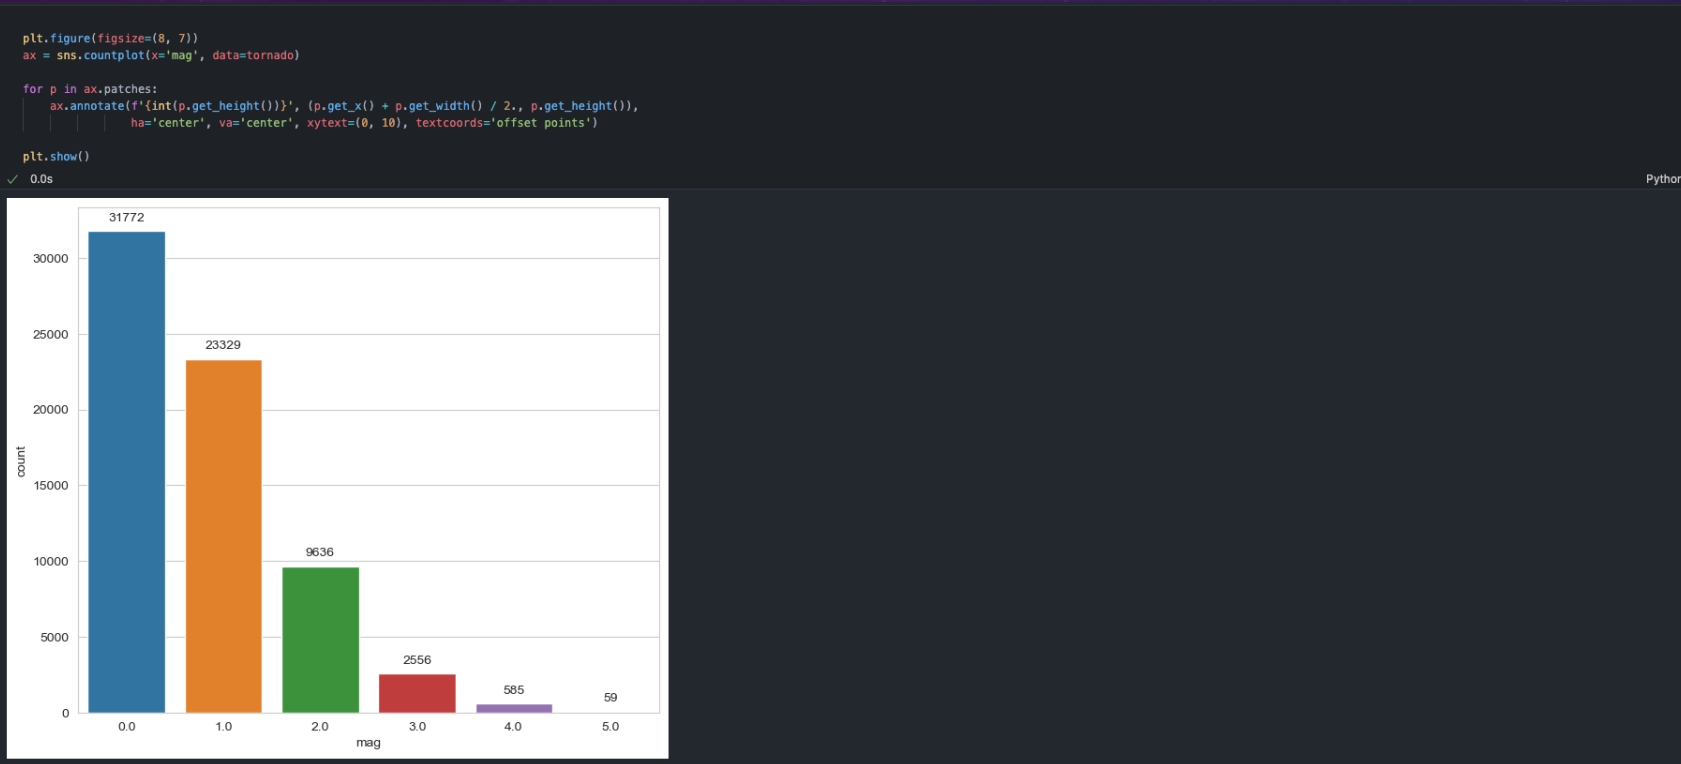
medianValue = tornado.loss.median()
stdValue = math.sqrt(49950)
new_data = np.random.normal(medianValue, stdValue, len(tornado[tornado['loss'].isnull()]))
tornado.loc[tornado['loss'].isnull(), 'loss'] = new_data
```

✓ 0.0s

Python



# Muestro cantidad de tornados para cada una de las magnitudes




# Muestro relación entre magnitud y fallecimientos

```
plt.figure(figsize=(7, 6))  
sns.scatterplot(x='mag', y='fat', data=tornado)  
plt.show()
```

✓ 0.1s

Python




# Muestra relación entre pérdidas materiales y magnitud

```
plt.figure(figsize=(6, 5))  
sns.scatterplot(x='mag', y='loss', data=tornado)  
plt.show()
```

✓ 0.1s

Python





# Compruebo cuántos tornados hay con un loss mayor a 1 billón

```
tornado[tornado.loss>1000000000]
```

✓ 0.0s

Python

	yr	mo	dy	date	time	st	stf	mag	inj	fat	...	elon	len	wid	ns	sn	f1	f2	f3	f4	fc
55272	2011	4	27	2011-04-27	14:05:00	1	1	5.0	145	72	...	-86.1511	132.00	2200	2	0	0	0	0	0	0
55288	2011	4	27	2011-04-27	15:43:00	1	1	4.0	1500	64	...	-86.7436	80.68	2600	1	1	63	125	73	0	0
55592	2011	5	22	2011-05-22	16:34:00	24	29	5.0	1150	158	...	-94.2213	21.62	1600	1	1	145	97	145	0	0
57584	2013	5	20	2013-05-20	13:56:00	36	40	5.0	212	24	...	-97.3999	13.85	1900	1	1	87	27	0	0	0
63647	2019	10	20	2019-10-20	19:58:00	44	48	3.0	0	0	...	-96.6833	15.76	1300	1	1	113	0	0	0	0
65427	2020	3	3	2020-03-03	00:32:00	43	47	3.0	220	5	...	-85.8904	60.13	1600	1	1	37	189	159	0	0

6 rows x 24 columns

Elimino estos 6 datos para que no causen errores en el modelo


```
tornado = tornado[tornado['loss'] <= 1000000000]
```

✓ 0.0s

Python



Muestro de nuevo el gráfico para comprobar correcto resultado





## Cambio tipo de los datos en las fechas

```
tornado['date'] = panda.to_datetime(tornado['date'])  
tornado['time'] = panda.to_datetime(tornado['time'], format='%H:%M:%S')
```

✓ 0.0s

Python

# Genero escalado

